



# Detection of Covid-19 Pneumonia via Chest X-Rays



**BCF 1**

Elangovan Karthikeyan  
Lam Wei Lin, Zoey  
Lee Pei Yee



## Current Situation ...

As COVID-19 becomes endemic, the **strain** on healthcare services worldwide **intensifies**, making it more difficult to **manage & detect** the serious complications of COVID-19.

We see the possibility of replacing the current, manual way of diagnosing pneumonia which is **unnecessarily exhausting limited resources** and **time-consuming** yet may not be highly effective due to the presence of **inevitable human errors**



# How can we improve the efficiency of diagnosing normal, viral and covid-induced pneumonia in healthcare settings?



Problem Definition

Data Preparation

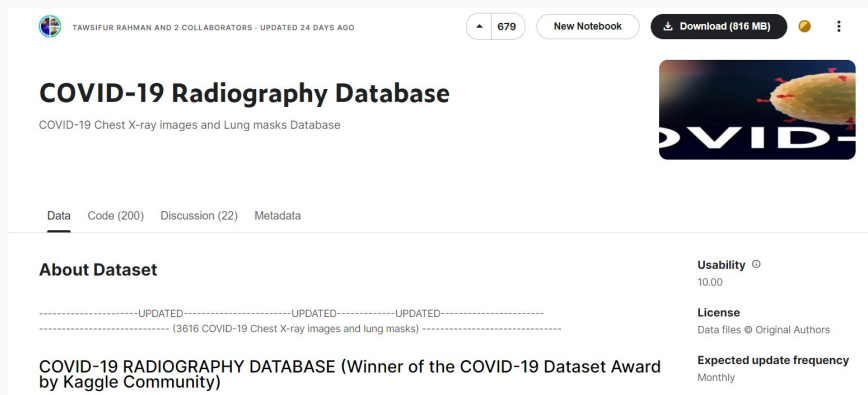
Exploratory Data Analysis

Machine learning

Insights

# PROBLEM DEFINITION...

## COVID-19 Chest X-Ray Image Dataset



- Found a dataset containing COVID-19 Chest X-Rays on Kaggle
- 21165 Chest X-Rays from 4 different categories
  - COVID\_19
  - Normal
  - Lung Opacity
  - Viral Pneumonia


# Our Solution...

To tackle this Multi-Class Classification Problem, we propose:

**Image Classification model to classify Chest X-Ray images into 3 classes (normal, viral pneumonia and COVID-induced pneumonia)**

## Approach:

- Clean and use only relevant data
- Split into train, validation and test set
- Apply Transfer Learning



# DATA PREPARATION...

```
▶ # Basic Libraries
import numpy as np
import pandas as pd
import seaborn as sb
import matplotlib.pyplot as plt # we only need pyplot
sb.set() # set the default Seaborn style for graphics
```

```
[2] #Testing for GPU
import tensorflow as tf
device_name = tf.test.gpu_device_name()
if device_name != '/device:GPU:0':
    raise SystemError('GPU device not found')
print('Found GPU at: {}'.format(device_name))
```

Found GPU at: /device:GPU:0

● Import Essential Libraries for analysis and machine learning such as Tensorflow (NEW!)

Importing Dataset using Kaggle API (kaggle.json file required)

```
[3] !pip install -q kaggle
from google.colab import files
files.upload()
! mkdir ~/.kaggle
! cp kaggle.json ~/.kaggle/
! chmod 600 ~/.kaggle/kaggle.json
! kaggle datasets list
```

Imported Dataset using Kaggle API

# DATA PREPARATION...

```
[16] files_normal = glob.glob('./COVID-19_Radiography_Dataset/Normal/images/*.png')
      files_COVID = glob.glob('./COVID-19_Radiography_Dataset/COVID/images/*.png')
      files_viral = glob.glob('./COVID-19_Radiography_Dataset/Viral_Pneumonia/images/*.png')
      files_Lung = glob.glob('./COVID-19_Radiography_Dataset/Lung_Opacity/images/*.png')
```

```
[17] print(len(files_normal))
      print(len(files_viral))
      print(len(files_COVID))
      print(len(files_Lung))
```

```
10192
1345
3616
6012
```

## Lung Opacity

- 6012 images in dataset
- Contains a mix of Viral Pneumonia and other lung infections

Overlapping data could cause accuracy of classifier to decrease

Other lung infections are not relevant to our goal

**Solution: Filter out the data**

# DATA PREPARATION...

Cleaning Data by Checking for Full White/Black Images in Files To Be Eliminated from dataset

```
[87] from PIL import Image

def checkFullWhiteOrBlack(folder):
    count = 0
    for i in range(len(folder)):
        img = Image.open(folder[i])
        clrs = img.getcolors()
        if len(clrs) == 1:
            count += 1
    return count

print("No Of Full White/Black Images in Files - Normal" , checkFullWhiteOrBlack(files_normal))
print("No Of Full White/Black Images in Files - Viral" , checkFullWhiteOrBlack(files_viral))
print("No Of Full White/Black Images in Files - COVID" , checkFullWhiteOrBlack(files_COVID))
```

```
No Of Full White/Black Images in Files - Normal 0
No Of Full White/Black Images in Files - Viral 0
No Of Full White/Black Images in Files - COVID 0
```

It can be observed that there are no full white/black images in all the files, hence no further cleaning is required.

## Checking for Blank Images


- There may be full black or white images within the files
- This would interfere with the accuracy of the classifier and the parameters used in it

## Solution: Check and filter out blank images

- But upon checking, we found no blank images
- Did not filter out any image



# DATA PREPARATION...



```
import numpy as np
np.random.seed(42)
files_normal = np.random.choice(files_normal, 1345, replace=False)
files_viral = np.random.choice(files_viral, 1345, replace=False)
files_COVID = np.random.choice(files_COVID, 1345, replace=False)
```

```
[18] print(len(files_normal))
      print(len(files_viral))
      print(len(files_COVID))
```

```
1345
1345
1345
```

## Data Imbalance

- 10192 Normal X-Rays
- 1345 Viral Pneumonia X-Rays
- 3616 COVID Pneumonia X-Rays

Highly Skewed to Normal Dataset

## Solution: Undersampling (NEW!)

- Use class with lowest number of datapoints as base
- Randomly pick data from other classes

Balances out the number of chest x-ray images for each category.

# Train Test Split ●●●

```
[ ] from sklearn.model_selection import train_test_split


normal_train, normal_test = train_test_split(files_normal, test_size=0.2, random_state=42, shuffle = True)
normal_valid, normal_test = train_test_split(normal_test, test_size=0.5, random_state=42, shuffle = True)
viral_train, viral_test = train_test_split(files_viral, test_size=0.2, random_state=42, shuffle = True)
viral_valid, viral_test = train_test_split(viral_test, test_size=0.5, random_state=42, shuffle = True)
COVID_train, COVID_test = train_test_split(files_COVID, test_size=0.2, random_state=42, shuffle = True)
COVID_valid, COVID_test = train_test_split(COVID_test, test_size=0.5, random_state=42, shuffle = True)
```

normal-train: 1076  
normal-valid: 134  
normal-test: 135  
Total Normal: 1345

viral-train: 1076  
viral-valid: 134  
viral-test: 135  
Total Viral: 1345

COVID-train: 1076  
COVID-valid: 134  
COVID-test: 135  
Total COVID: 1345

Total Train dataset: 3228  
Total Valid dataset: 402  
Total Test dataset: 405



# Encoding Classes in Data ...

```
✓ [23] #Assigned class for normal - 0, viral - 1, COVID - 2 - check if we wanna assign 0 to viral as well!  
0s  
  
df_train_normal['class'] = pd.Series([0 for x in range(len(df_train_normal.index))], index=df_train_normal.index)  
df_train_viral['class'] = pd.Series([1 for x in range(len(df_train_viral.index))], index=df_train_viral.index)  
df_train_COVID['class'] = pd.Series([2 for x in range(len(df_train_COVID.index))], index=df_train_COVID.index)  
  
df_test_normal['class'] = pd.Series([0 for x in range(len(df_test_normal.index))], index=df_test_normal.index)  
df_test_viral['class'] = pd.Series([1 for x in range(len(df_test_viral.index))], index=df_test_viral.index)  
df_test_COVID['class'] = pd.Series([2 for x in range(len(df_test_COVID.index))], index=df_test_COVID.index)  
  
df_valid_normal['class'] = pd.Series([0 for x in range(len(df_valid_normal.index))], index=df_valid_normal.index)  
df_valid_viral['class'] = pd.Series([1 for x in range(len(df_valid_viral.index))], index=df_valid_viral.index)  
df_valid_COVID['class'] = pd.Series([2 for x in range(len(df_valid_COVID.index))], index=df_valid_COVID.index)  
df_train_normal.head()
```

	filename	class
0	./COVID-19_Radiography_Dataset/Normal/images/N...	0
1	./COVID-19_Radiography_Dataset/Normal/images/N...	0
2	./COVID-19_Radiography_Dataset/Normal/images/N...	0
3	./COVID-19_Radiography_Dataset/Normal/images/N...	0
4	./COVID-19_Radiography_Dataset/Normal/images/N...	0

0 - Normal

1 - Viral Pneumonia

2 - COVID Pneumonia

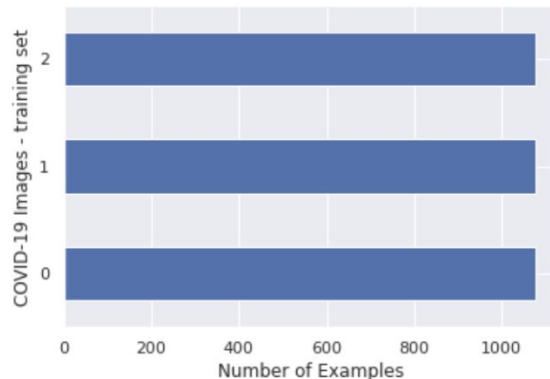
# Exploratory Data Analysis ...

Inspecting the distribution of classes in the train dataset

```
print(df['class'].value_counts().sort_index())
```

```
ax = df['class'].value_counts().sort_index().plot.barh()  
ax.set_xlabel("Number of Examples", fontsize=12)  
ax.set_ylabel("COVID-19 Images - training set", fontsize=12)  
plt.show()
```

```
0    1076  
1    1076  
2    1076  
Name: class, dtype: int64
```




# Exploratory Data Analysis ...

## Raw Observations Of Image Data

```
# randomly select 3 of each
select_norm = np.random.choice(train_normal, 3, replace = False)
select_viral = np.random.choice(train_viral, 3, replace = False)
select_covid = np.random.choice(train_covid, 3, replace = False)

# plotting 3 x 3 image matrix
fig = plt.figure(figsize = (8,6))
for i in range(9):
    if i < 3:
        fp = select_norm[i]
        label = 'NORMAL'
    elif i < 6:
        fp = select_viral[i - 3]
        label = 'VIRAL PNEUMONIA'
    else:
        fp = select_covid[i - 6]
        label = 'COVID PNEUMONIA'
    ax = fig.add_subplot(3, 3, i+1)

    # to plot without rescaling, remove target_size
    fn = image.load_img(fp, target_size = (100,100), color_mode='grayscale')
    plt.imshow(fn, cmap='Greys_r')
    plt.title(label)
    plt.axis('off')
plt.show()
```



# Exploratory Data Analysis ...

## Average Comparison Of Image Data (NEW!)

Converting Images to Matrix



Finding the average image for each class

```
def img2np(list_of_filename, size = (64, 64)):
    # iterating through each file
    for fn in list_of_filename:
        fp = fn
        current_image = image.load_img(fp, target_size = size,
                                         color_mode = 'grayscale')


        # covert image to a matrix
        img_ts = image.img_to_array(current_image)
        # turn that into a vector / 1D array
        img_ts = [img_ts.ravel()]
        try:
            # concatenate different images
            full_mat = np.concatenate((full_mat, img_ts))
        except UnboundLocalError:
            # if not assigned yet, assign one
            full_mat = img_ts
    return full_mat

# run it on our folders
normal_images = img2np(train_normal)
viral_images = img2np(train_viral)
covid_images = img2np(train_covid)
```


```
def find_mean_img(full_mat, title, size = (64, 64)):
    # calculate the average
    mean_img = np.mean(full_mat, axis = 0)
    # reshape it back to a matrix
    mean_img = mean_img.reshape(size)
    plt.imshow(mean_img, vmin=0, vmax=255, cmap='Greys_r')
    plt.title(f'Average {title}')
    plt.axis('off')
    plt.show()
    return mean_img

norm_mean = find_mean_img(normal_images, 'NORMAL')
viral_mean = find_mean_img(viral_images, 'VIRAL PNEUMONIA')
covid_mean = find_mean_img(covid_images, 'COVID PNEUMONIA')
```


Average NORMAL



Average VIRAL PNEUMONIA



Average COVID PNEUMONIA



# Exploratory Data Analysis ...


Contrast between Average Images Of  
Normal, Viral Pneumonia and Covid  
Pneumonia(NEW!)

```
[38] fig = plt.figure(figsize = (18,6))
      for i in range(3):
          if i == 0:
              contrast_mean = norm_mean - viral_mean
              label = 'Difference Between Normal & Viral Pneumonia Average'
          elif i == 1:
              contrast_mean = norm_mean - covid_mean
              label = 'Difference Between Normal & Covid Pneumonia Average'
          else:
              contrast_mean = viral_mean - covid_mean
              label = 'Difference Between Viral & Covid Pneumonia Average'


          ax = fig.add_subplot(1, 3, i+1)

          plt.imshow(contrast_mean, cmap='bwr',)
          plt.title(label)
          plt.axis('off')
      plt.show()
```


Difference Between Normal & Covid Pneumonia Average



Difference Between Normal & Viral Pneumonia Average



Difference Between Viral & Covid Pneumonia Average



# Exploratory Data Analysis ...

## Eigenimages from Principal Component Analysis (PCA) (NEW!)




### Dimension reduction technique — Principal Component Analysis (PCA)

- Visualize the components that best describes each class.
- Eigenimages — Plotted from matrix shaped by components from PCA
  - Principal components that describe **70%** of variability for each class.



# Exploratory Data Analysis

## Image Value Distribution (NEW!)



3 clusters — Normal, Viral, Covid

### Normal:

- Mean - 100 to 150
- Standard Deviation - 60 to 70

### Viral:

- Mean - 100 to 150
- Standard Deviation - 50 to 80

### Covid:


- Mean - 125 to 175
- Standard Deviation - 40 to 60

# Exploratory Data Analysis ...

Image Value Distribution (NEW!)


black -  $\text{rgb}(0,0,0)$

White -  $\text{rgb}(255,255,255)$




## Normal:

- 100 to 150
- Majority: 120 to 125



## Viral:

- 105 to 145
- Majority: 130 to 135



## Covid:

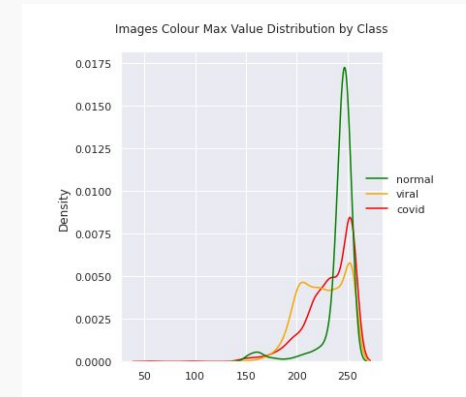
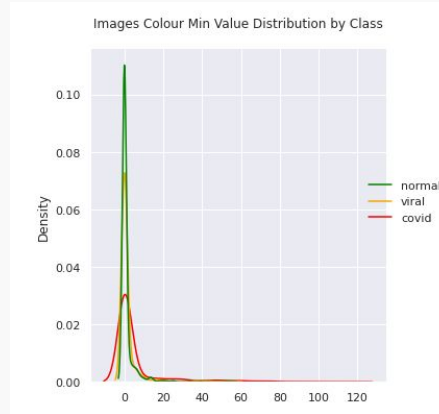

- 125 to 175
- Majority: 145 to 150

# Exploratory Data Analysis ...

## Image Value Distribution — Pixel Density plots (NEW!)

black - `rgb(0,0,0)`

White - `rgb(255,255,255)`



### Mean Values:

- Normal: 120
- Viral: 130
- Covid: 145

### Min Values:

- Normal: 0
- Viral: 0
- Covid: 0

### Max Values:

- Normal: 245
- Viral: 200 to 250
- Covid: 230 to 250

# Machine Learning ...

## Preparation of Pipeline to Load images



```
▶ # convert the dataframe into 2 lists to use for filename and labels
train_filenames_list = df["filename"].tolist()
train_labels_list = df["class"].astype('int32').tolist()

# convert the dataframe into 2 lists to use for filename and labels
val_filenames_list = df_val["filename"].tolist()
val_labels_list = df_val["class"].astype('int32').tolist()

# convert the dataframe into 2 lists to use for filename and labels
test_filenames_list = df_test["filename"].tolist()
test_labels_list = df_test["class"].astype('int32').tolist()

#number of classes
num_classes = 3
```

```
# Reads an image from a file, decodes it into a tensor, and resizes it
# to a fixed shape.
```

```
img_rows, img_cols = 299,299
```

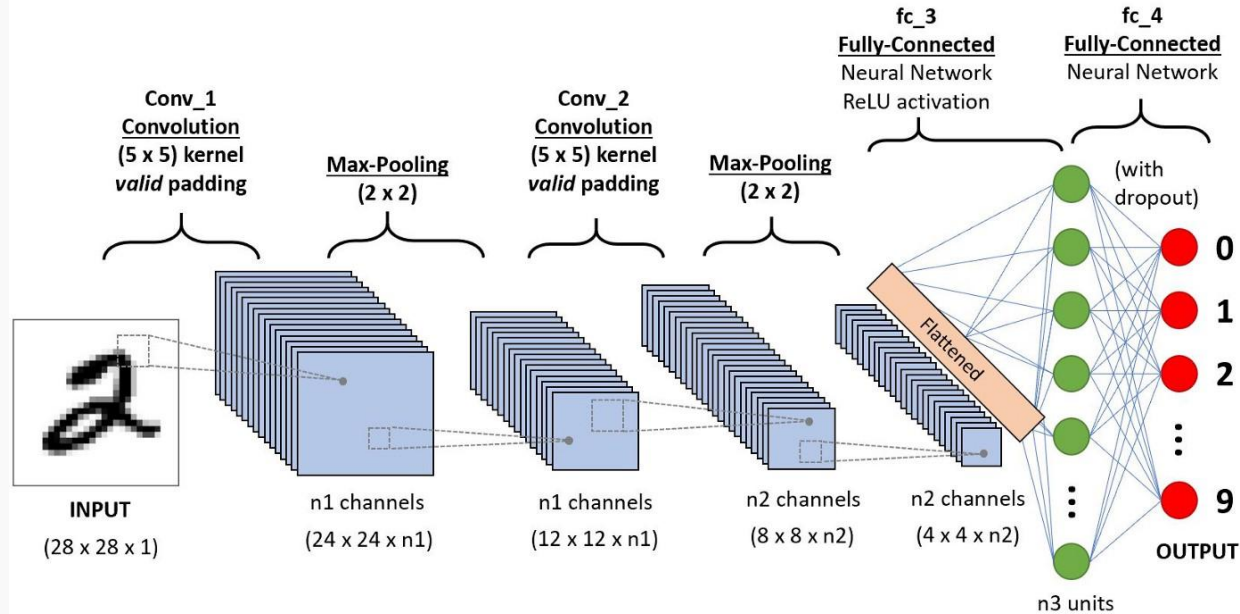
```
def _parse_function(filename, label):
    image_string = tf.io.read_file(filename)
    image_decoded = tf.image.decode_jpeg(image_string,channels=3)
    image_resized = tf.image.resize(image_decoded, [img_rows, img_cols])
    label = tf.one_hot(label, num_classes)
    return image_resized, label
```

```
[46] train_dataset = tf.data.Dataset.from_tensor_slices((filenames, labels))
      train_dataset = train_dataset.map(_parse_function)
      train_dataset = train_dataset.repeat(100)
      train_dataset = train_dataset.batch(32)
```

```
[47] valid_dataset = tf.data.Dataset.from_tensor_slices((val_filenames, val_labels))
      valid_dataset = valid_dataset.map(_parse_function)
      valid_dataset = valid_dataset.repeat(100)
      valid_dataset = valid_dataset.batch(32)
```

```
[48] test_dataset = tf.data.Dataset.from_tensor_slices((test_filenames, test_labels))
      test_dataset = test_dataset.map(_parse_function)
      test_dataset = test_dataset.repeat(100)
      test_dataset = test_dataset.batch(32)
```

# Convolutional Neural Networks (CNN) (NEW!)



# Convolutional Neural Networks (CNN) (NEW!)

```
[55] from tensorflow.keras import layers, models

cnn_model = tf.keras.models.Sequential()
cnn_model.add(layers.BatchNormalization(input_shape=(299,299,3)))
cnn_model.add(layers.Conv2D(filters = 128, kernel_size = (3, 3), activation = 'relu'))
cnn_model.add(layers.MaxPooling2D((2, 2)))
cnn_model.add(layers.Dropout(0.3))

cnn_model.add(layers.Conv2D(filters = 64, kernel_size = (3, 3), activation = 'relu'))
cnn_model.add(layers.MaxPooling2D((2, 2)))
cnn_model.add(layers.Dropout(0.35))

cnn_model.add(layers.Conv2D(filters = 32, kernel_size = (3, 3), activation = 'relu'))
cnn_model.add(layers.Flatten())
cnn_model.add(layers.Dense(units = 128, activation = 'relu'))
cnn_model.add(layers.Dropout(0.3))

cnn_model.add(layers.Dense(units = 3, activation = 'softmax'))

cnn_model.compile(optimizer = 'adam',
                  loss = 'categorical_crossentropy',
                  metrics = ['accuracy'])


cnn_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
batch_normalization (Batch Normalization)	(None, 299, 299, 3)	12
conv2d (Conv2D)	(None, 297, 297, 128)	3584
max_pooling2d (MaxPooling2D)	(None, 148, 148, 128)	0
dropout (Dropout)	(None, 148, 148, 128)	0
conv2d_1 (Conv2D)	(None, 146, 146, 64)	73792
max_pooling2d_1 (MaxPooling2D)	(None, 73, 73, 64)	0
dropout_1 (Dropout)	(None, 73, 73, 64)	0
conv2d_2 (Conv2D)	(None, 71, 71, 32)	18464
flatten (Flatten)	(None, 161312)	0
dense (Dense)	(None, 128)	20648064
dropout_2 (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 3)	387



=====

Total params: 20,744,303  
Trainable params: 20,744,297  
Non-trainable params: 6



# Convolutional Neural Networks (CNN) (NEW!)

```
[59] plot_loss_acc(history_cnn)
```



```
[60] metrics = cnn_model.evaluate(test_dataset, steps=50)
print("model accuracy:", metrics[1])
```

```
50/50 [=====] - 7s 144ms/step - loss: 0.3920 - accuracy: 0.8475
model accuracy: 0.8475000262260437
```





# Transfer Learning (NEW!)

## Traditional ML


vs

## Transfer Learning

- Isolated, single task learning:
  - Knowledge is not retained or accumulated. Learning is performed w.o. considering past learned knowledge in other tasks



- Learning of a new tasks relies on the previous learned tasks:
  - Learning process can be faster, more accurate and/or need less training data




## Key Strengths Over Traditional ML models

- Don't require training from **scratch**
- **Computationally less expensive** and efficient to achieve better results
- Achieve **optimal performance faster** as they leverage knowledge from previously trained models that already understand the features



# Transfer Learning (NEW!)

Pre-trained model Used : **Inception V3**



## Key aspects:

- Good accuracy
- Low parameters

Feasible to deploy this **lightweight** model in **healthcare** settings

# Transfer Learning ...

## 1. Creating a base model from InceptionV3

```
[61] # InceptionV3 Model - trained on image net
      base_model = InceptionV3(include_top=False, weights='imagenet')

      Downloading data from https://storage.googleapis.com/tensorflow/k87916544/87910968 [=====] - 1s 0us/step
      87924736/87910968 [=====] - 1s 0us/step
```

```
=====
Total params: 21,802,784
Trainable params: 21,768,352
Non-trainable params: 34,432
=====
```

## 2. Creating a simple classification head so that we can train it

```
#add a global spatial average pooling layer
x = base_model.output

x = GlobalAveragePooling2D()(x)

# let's add a fully-connected layer
x = Dense(1024, activation='relu')(x)

# 3 Classes for Detection of Pneumonia
num_classes = 3
predictions = Dense(num_classes, activation='softmax')(x)
```

# Transfer Learning ...

## 3. Freezing convolutional base layers except BatchNorm, to use as a feature extractor

```
# Creating the model that we will train
model = Model(inputs=base_model.input, outputs=predictions)

# Freeze the InceptionV3 layers except the BatchNorm
for layer in base_model.layers:
    if 'batch' in layer.name:
        print(layer.name)
        layer.trainable = True
    else:
        layer.trainable = False
```

## 4. Selecting an Optimizer - The Adam Optimization algorithm

```
# tf.keras optimizer
opt = tf.keras.optimizers.Adam(lr)
```

## 5. Compiling The Model with Loss Function

```
# compile the model
model.compile(optimizer= opt, loss='categorical_crossentropy', metrics=[ 'accuracy' ])
```

# Transfer Learning ...

## 6. Creating Checkpoints To Save Best Weights

```
#Creating Checkpoints to save best weights  
checkpoint2 = ModelCheckpoint('./checkpoints/best_weights_InceptionV3Model.hdf5', verbose=1, save_best_only=True, mode='auto')
```

## 7. Setting up some hyperparameters

```
# Set up some Hyperparameters
```

```
batch_size = 32  
lr = 0.001
```

```
training_data_size = len(normal_train) + len(viral_train)+len(COVID_train)  
  
train_steps = int(training_data_size/batch_size) #training data / batch size = 3228/32  
print(train_steps)  
val_steps = 50  
epochs = 5
```

```
100
```

# Transfer Learning ...

## Fitting the model

```
history = model.fit( train_dataset, steps_per_epoch = train_steps,
                    epochs = epochs,
                    validation_data = valid_dataset,
                    validation_steps = val_steps,
                    callbacks=[checkpoint2])
```

```
Epoch 1/5
100/100 [=====] - ETA: 0s - loss: 0.6543 - accuracy: 0.7584
Epoch 1: val_loss improved from inf to 0.91719, saving model to ./checkpoints/best_weights_InceptionV3Model.hdf5
100/100 [=====] - 55s 463ms/step - loss: 0.6543 - accuracy: 0.7584 - val_loss: 0.9172 - val_accuracy: 0.6569
Epoch 2/5
100/100 [=====] - ETA: 0s - loss: 0.2919 - accuracy: 0.8844
Epoch 2: val_loss improved from 0.91719 to 0.35303, saving model to ./checkpoints/best_weights_InceptionV3Model.hdf5
100/100 [=====] - 45s 447ms/step - loss: 0.2919 - accuracy: 0.8844 - val_loss: 0.3530 - val_accuracy: 0.8581
Epoch 3/5
100/100 [=====] - ETA: 0s - loss: 0.2241 - accuracy: 0.9134
Epoch 3: val_loss improved from 0.35303 to 0.18712, saving model to ./checkpoints/best_weights_InceptionV3Model.hdf5
100/100 [=====] - 45s 447ms/step - loss: 0.2241 - accuracy: 0.9134 - val_loss: 0.1871 - val_accuracy: 0.9125
Epoch 4/5
100/100 [=====] - ETA: 0s - loss: 0.1654 - accuracy: 0.9375
Epoch 4: val_loss did not improve from 0.18712
100/100 [=====] - 44s 439ms/step - loss: 0.1654 - accuracy: 0.9375 - val_loss: 0.2122 - val_accuracy: 0.9050
Epoch 5/5
100/100 [=====] - ETA: 0s - loss: 0.1099 - accuracy: 0.9603
Epoch 5: val_loss did not improve from 0.18712
100/100 [=====] - 44s 438ms/step - loss: 0.1099 - accuracy: 0.9603 - val_loss: 0.2306 - val_accuracy: 0.9050
```

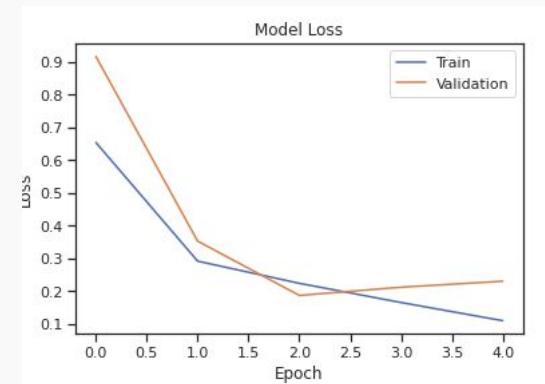

# Transfer Learning ...

## Visualising Loss Over Epochs

```
#Visualising model loss over epochs

def plot_train(hist):
    h = hist.history
    if 'acc' in h:
        meas='acc'
        loc='lower right'
    else:
        meas='loss'
        loc='upper right'
    plt.plot(hist.history[meas])
    plt.plot(hist.history['val_'+meas])
    plt.title('model '+meas)
    plt.ylabel(meas)
    plt.xlabel('epoch')
    plt.legend(['train', 'validation'], loc=loc)
```

**Increase in Validation Accuracy**  
**Drop in Validation Loss**





# Transfer Learning ...

Further training for more epochs using more callbacks for optimisation in training process

```
reduce_LR = tf.keras.callbacks.ReduceLROnPlateau(monitor='val_loss', factor = 0.9, patience=2, cooldown=1, min_lr = 0.00001)
lr_print = tf.keras.callbacks.LambdaCallback(on_epoch_begin=lambda epoch, logs: print("lr:", K.eval(model.optimizer.lr)))
early_stopping = tf.keras.callbacks.EarlyStopping(monitor = 'val_loss', mode = 'min', verbose = 1, patience = 4)
```

```
## Train the model with validation
history2 = model.fit ( train_dataset, steps_per_epoch = train_steps,
                      epochs = 10,
                      validation_data = valid_dataset,
                      validation_steps = val_steps,
                      callbacks=[checkpoint2, lr_print, reduce_LR, early_stopping])
```



**Some improvement in validation accuracy**

- **Some signs of overfitting**

# Outcomes

## Evaluating model with best weights:

- Test Accuracy: 0.9350000023841858
- Outstanding accuracy of 93.5%.


## Confusion Matrix:

- Our model has a **significantly high true positives** seen from 0.91(normal), 0.99(viral) and 0.90(COVID)
- There is a relatively significant chance of classifying **viral pneumonia** (0.01) and **COVID X-Rays** (0.08) as **normal** (0.09)
- The chances of wrongly classifying **COVID as viral pneumonia** (0.02) is **higher** than classifying **viral pneumonia as COVID** (0.00)

```
[79] model.load_weights('./checkpoints/best_weights_InceptionV3Model.hdf5')
```

```
[80] metrics = model.evaluate(test_dataset, steps=50)  
print("Test accuracy:", metrics[1])
```

```
50/50 [=====] - 9s 178ms/step - loss: 0.1605 - accuracy: 0.9350  
Test accuracy: 0.9350000023841858
```





# Outcomes ...

## Solving our problem:

Our model can be used in real life scenarios with its high accuracy, and scalability, and can be easily deployed onto web applications for healthcare professionals to use to accurately distinguish between Normal, Viral Pneumonia and COVID Cases, enabling faster diagnosis and treatment of viral and COVID-19 cases.

In the future, we can easily modify our model for other medical use cases (e.g detection of other lung infections like tuberculosis )

# Improvements and further Modifications ...

## Improving dataset:

- Image Augmentation — Genetic Algorithm, Keras ImageDataGenerator
- Adopt oversampling
- Adding on more X-Ray images to classes with less data

## Other Pre-trained models:

- VGG16, ResNet

## Other factors:

- Symptoms of COVID and viral pneumonia
- Other forms of Medical Diagnosis Images (e.g Brain CT scans)

# What Have We Learned ...

## Exploratory Data Analysis

- Pixel Matrix analysis
- Eigenimages from Principal Component Analysis (PCA)
- Image Value Distributions

## Machine Learning

- Convolutional Neural Networks (CNN)
- Transfer Learning
  - Tensorflow and Keras
  - Pre-trained Model - InceptionV3
  - Optimisers - Adam Optimisation Algorithm
  - Checkpoints, Callbacks



# THANK YOU



**Stay safe &  
remember to mask up indoors!**